Operating System for the Raspberry Pi

A Kernel and simple Shell for the 64-bit Raspberry Pi 3, developed from an historical perspective on the structure of early Operating Systems.

Sam Whitehead - 14325283
`psysrw@nottingham.ac.uk`
Msci Computer Science

COMP4029: Individual Programming Project
Project Supervisor: Steve Bagley
University of Nottingham

Abstract

In this project I aimed to create an Operating System (OS) for the Raspberry Pi (RPi) 3. I originally intended to follow the design of the Unix system from the beginning, but at the start my work did not seem to be heading in that direction.

Over the course of developing this OS, I learned a lot about which OS components are needed at which stages of development. This project has followed a logical "first discovery" journey into OS development, which shadowed the historic major advances in OS technology.

The OS I have created provides a kernel, with a system call interface, most of the infrastructure for a multiple-process execution model, an implementation of the exFAT Filesystem (FS), and a basic shell which allows users to run several different basic commands.

The final OS is not intended for daily use, but rather it is an example of a typical project for those wanting to learn about OS development. The source code for the project is made freely available online as a learning tool.

Contents

1 Introduction

First, an introduction to some of the concepts on which this project relies.

1.1 What is ARM64?

- ARM64 (or more commonly Aarch64) is the name given to the 64 bit architecture of arm Central Processing Units (CPUs).

1.2 What is an Operating System? (And Unix)

Operating Systems (OSes) are programs which run directly on top of computer hardware and allow other software to run on the computer. They come in many different varieties, but this project will be based on the "microkernel" design. Unix was an OS developed at Bell Labs of AT&T in the 1970s. It was originally written in assembly language for the specific target machine, the PDP7, but for its 4th version it was rewritten in C, a new language at the time (also created at Bell Labs). Being written in C made the OS very portable, as only the compiler had to be ported for the entire system to run on a new machine. This made Unix very popular, and soon there were many Unix-like OSes. 'A Unix-like OS is one that behaves in a similar manner to a Unix system' ([21]). There are many examples of Unix-like OSes, including some very popular ones like Linux, MacOS, and the family of OSes known as the BSDs. The OS I am creating in this project aims to be Unix-like in its behaviour.

1.3 Why the Raspberry Pi?

The Raspberry Pi (RPi) is a cheap Single Board Computer (SBC) powered by an arm System-On-a-Chip (SOC). This project focuses on the RPi 3, the first iteration of the RPi to use a 64-bit CPU. The RPi is ideal for a project like this, as it is complex enough that creating an OS for it is a challenge, but also modern enough that no out-of-date technologies will be necessary to work with it. The RPi board includes a number of different components that would require drivers to function, so there would be plenty of room for additional features if the project moved faster than anticipated.

The RPi is popular among hobbyist OS developers, and its hardware schematics and technical manuals are made available, so there are plenty of resources available online for documentation of technical details. The OS I have developed is intended to be a learning resource for others who wish to discover more about OS design. The source code will be made available under an open source license.

2 Motivation

My personal motivation for choosing to create an Operating System (OS) for this project stems from my interest in low-level programming. I have an interest in understanding how computers work at the lowest level, and the best way to learn how an OS is implemented and the reasons for design decisions will be to implement my own OS at that low level.

With this project, I wanted to develop a deeper understanding of the inner workings of OSes and how the programs I run on my computer interact with the hardware they run on. I wanted to understand the entire software stack, and one large gap in my knowledge was the internals of the Operating System. It is a black box to most students of Computer Science, and I was determined to fill that gap in my own understanding. I wanted to know how system calls work, and how they relate to the concept of a "secure monitor". I already knew how filesystems work in theory, but I wanted to experience implementing one to gain a practical appreciation of the topic.

This project is needed because although OSes are a common hobby project, there are no OS development projects which have aimed to highlight the motivation behind the choices that have been made by the largest OSes. No project exists which shows clearly the reasons behind why OS development as a field has taken the direction that it has. Why does every major OS use a software interrupt based system call architecture to distinguish between user-level code and privileged kernel operations? What are some of the reasons for choosing to include hardware drivers in the kernel of an OS instead of running those drivers as user code? These questions, and others like them, have not been answered to a suitable level by other projects.

3 Related work

3.1 Filesystem code and memory management

The project "rpi-boot" [Rw6] contains code for a File Allocation Table (FAT) Filesystem (FS) and also some code for common libc functions. It includes an implementation of the `malloc` function [Rw1], used to allocate memory.

I will use the FS code from this source as inspiration for the FS of my Operating System (OS). The `malloc` implementation will also be used as the basis for my memory management system.

3.2 Central Processing Unit (CPU) mode initialisation for arm CPUs

A project [Rw3] which contains assembly code to switch arm 64-bit architecture CPUs into the correct operation mode for an OS. This code will be very useful when I write the boot code, as I will want to setup the processor into the correct Exception Level (EL) for the kernel.

This project also contains example code for Exception handlers, which is something I will need in order to implement System Calls (subsubsection 8.2.2).

3.3 Universal Asynchronous Receiver-Transmitter (UART) serial Input / Output (IO)

A set of tutorials on Github [Rw4] contains code which enables debugging via the UART serial pins on the Raspberry Pi (RPi). I will use this code as a reference when I implement a library for the two UART outputs for the RPi.

This project's later lessons also implement code for the framebuffer and for arm ELs, but I will instead reference documentation and other implementations when I work on these for my OS.

3.4 Other related works

Other related works are different OSes. The Linux kernel [Rw7] is a well known monolithic kernel, and supports numerous different CPU architectures, including Aarch64. I will look at the Linux kernel's implementation of system calls when I start working on the system call interface for my OS.

Another OS which I will look at is NetBSD [Rw2] and the rest of the BSDs. The source code for NetBSD will be simpler and easier to understand than that of Linux, so I will reference BSD implementations of any OS components which are overly complicated in the Linux kernel.

The RiscOS operating System [Rw5] was the first OS written for arm CPUs. It has been kept updated and a version is available which runs on the RPi. The entire OS is written in arm assembly language, so it will not be particularly useful to me, but I can take a look at its implementation of any components I get particularly stuck on.

References for the Related work

[Rw1]   Doug Lea. *DL malloc, a version of `malloc/realloc/free`*. Version 2.8.6. URL: http://gee.cs. oswego.edu/pub/misc/malloc.c (visited on 2021-11-28).

[Rw2]   *NetBSD*. URL: https://github.com/NetBSD/src (visited on 2021-11-28).

[Rw3]   *raspberry-pi-os: Learning OS development using Linux kernel and RPi*. URL: https://github. com/s-matyukevich/raspberry-pi-os (visited on 2021-11-28).

[Rw4]   *raspi3-tutorial: Bare metal RPi 3 tutorials*. URL: https://github.com/bztsrc/raspi3-tutorial (visited on 2021-11-28).

[Rw5]   *RiscOS*. Version 5.29. URL: https://www.riscosopen.org/content/ (visited on 2021-11-28).

[Rw6]   *rpi-boot: A second stage bootloader for the RPi*. URL: https://github.com/jncronin/rpi-boot (visited on 2021-11-28).

[Rw7]   *The Linux kernel*. Version v5.15. URL: https://git.kernel.org/pub/scm/linux/kernel/git/ torvalds/linux.git/snapshot/linux-5.15.tar.gz (visited on 2021-11-28).

Task	Duration (weeks)
Kernel	
Bootloader	1
Graphics driver	2
Syscall interface	4
Memory allocation	2
SD card driver	2
Filesystem	9
Shell	
Debug terminal	2
Graphical console	3
Kernel shell interface	4
Basic shell commands	2
Tools and libraries	
Setup build environment	1
IO and strings libraries	2
libc	8
Documents	
Dissertation	7

Table 1: The plan of work for the project, including how many weeks I think each component will take.

4 Description of the work

This project will produce a piece of software. The software will aim to achieve two main goals: to be an Operating System (OS) and to be a useful starting point for anyone wanting to teach themselves OS development.

4.1 The goals of the project

To meet the first goal, the software produced should qualify as an OS. For this, it should have its code organised into a "kernel" structure. There should be drivers for the essential hardware, such as the serial port (for a debug console), the framebuffer (for graphical output), and the eMMC chip, which controls access to SD cards on the Raspberry Pi (RPi). There should also be enough filesystem code to allow for reading and writing of files, and there should be implementations of familiar standard library file I/O wrapper functions, including `fread` and `fwrite`.

For the second goal (being a useful starting point for someone learning about OS development), the OS needs to be simple enough for a competent programmer to get a basic understanding of how all of the components fit together, within only a couple of hours of first studying the source code. As with the first goal, it will need to include all of the required features for a basic OS, and these were described previously. It should also have plenty of room for future development, as I believe that one of the best ways to understand a topic (and to become familiar with an existing project) is to try to develop a new feature involving the topic or extend the project.

4.2 Specification of the project

I specified the components of the project, and these components are described below. The components are split into two parts: the Tasks of the OS, and a selection of Work Packages I will work on.

4.2.1 Tasks

Tasks are the individual components of the project, separated into the components of the Kernel and Shell, the miscellaneous parts, and the documents I will need to produce for the project. These tasks are enumerated in Table 1 on page 6 and then assembled into a Gantt chart in Figure 1.

Some of the tasks are different to the ones in the original table from the project proposal, which can be found in the Appendix (Table 3). Some tasks have changed places, and some have been given extra time. Also, some of the tasks have been completely removed as I decided that I will not have enough time to complete them (primarily: init system and service manager, and multithreading).

Work package	Description of the deliverable
Debug terminal using the serial port	A debug terminal which is printed to the RPi's serial port. The debug terminal should print important values while the OS is starting up, allowing basic debugging by verifying that memory addresses are set correctly and that functions are performing their behaviours as expected.
A basic kernel which can interact with the hardware	A program that is invoked from the bootloader and contains drivers for some essential hardware such as the eMMC controller and the framebuffer.
A shell built into the kernel	A simple shell which takes user input one line at a time and executes commands based on the first word of the input line. This shell will run as part of the kernel and will execute at the same privilege level as the rest of kernel code.
A separate shell and commands	A version of the shell that is packaged as a separate executable compared with the kernel. The commands that the shell executes will also be separate executables stored in the FS. The kernel will execute the shell using a binary file loader, and the shell will contain the same loader for executing commands.

Table 2: The work packages of the project and their deliverables.

4.2.2 Work Packages

A *Work Package* is a collection of tasks with a well-defined deliverable. Not every task will be part of a work package, and some tasks will be work packages on their own. Table 2 contains all of the work packages for this project.

5 Methodology

5.1 Setting up the build system

For this project, I could not use a traditional compilation setup. Normally, a program runs in a hosted environment, which means that it can use a standard library for some basic functions. In the case of C code, this is `libc`. My Operating System (OS) runs on *bare metal*, so it cannot use any of the standard header files provided by `libc`, which means that the entire C standard library is unavailable. I still have access to some header files, including `stddef.h` and `stdint.h`, as these are provided as part of the compiler. There are other header files available in this *freestanding* environment, but these two were the main ones for the first stages of development, since they define types such as `uint32_t` and `size_t`.

In addition to the compiler flags needed for freestanding development, I also needed a completely different compiler. This relates to the *Architecture* of the system running the compiler. The Raspberry Pi I used for this project has an architecture of `Aarch64`, or 64-bit ARM, whereas desktop PCs are usually `x86` or `x86_64` architecture computers. This means I needed a compiler that generates machine code for `Aarch64` based systems. Such a compiler would be called a *Cross-compiler* if it runs on a different architecture compared with the machine code it produces.

5.1.1 Cross-compilation terminology and tools

- **Build**: The architecture of the computer used to compile the program.

- **Host**: The architecture of the computer the program will run on.

- **Target**: The architecture the program builds for (only used when the program is or contains a compiler).

For this project I used the LLVM suite of tools, as it was designed to be a cross-compiler first, so all of the tools (compiler, assembler, binary manipulation tools, etc.) are capable of handling the required Aarch64 host system. The compiler, `clang`, takes a command-line argument which indicates the architecture of the host system (LLVM ignores the difference between host and target, calling them both 'target').

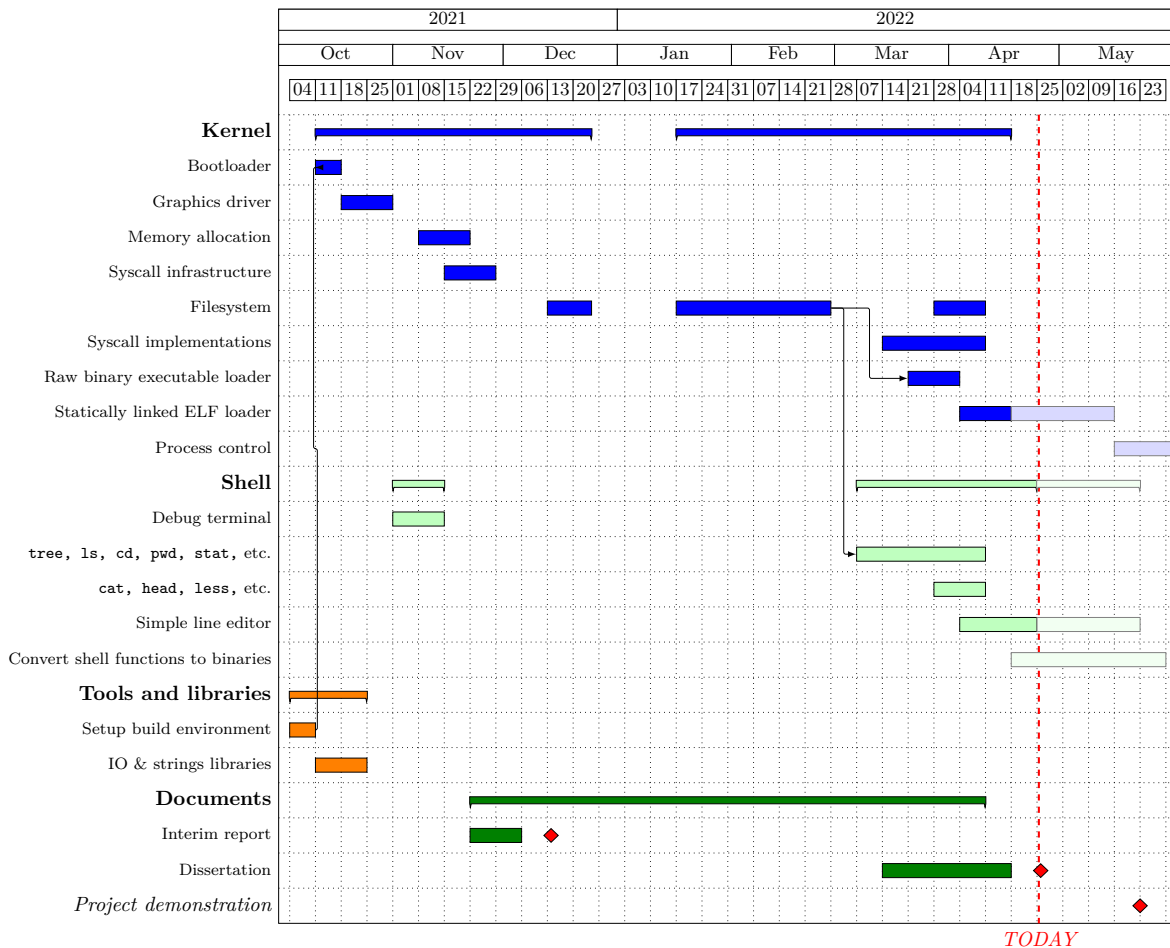### 5.2 The build process - Make and Makefile

TODO

Figure 1: Gantt chart

5.3 Software tools used to aid development

5.3.1 Source code editing

For editing the source code of the OS, I use the text editor Vim [22] with several plugins to assist me in writing correct code. I am using the CoC extension [7] for autocompletion, and the clangd language server [5] extension for CoC in order to provide source code linting for C source files.

For source code formatting, I run the program `clang-format` [6] on the C source and header files. I have defined a style guide for the project in a `.clang-format` configuration file, which tells `clang-format` how to format my code. This results in a consistent style throughout the entire project, causing the source code to be easier to read and understand.

5.3.2 Raspberry Pi 3 Emulator

QEMU [1] is an emulator with out-of-the-box support for emulating the host architecture. This means that on an `x86` machine, we can emulate another `x86` machine. This isn't enough for emulating a Raspberry Pi (RPi), as the RPi uses the Aarch64 architecture. Thankfully, there is a package on most linux distributions, ususally called something like `qemu-arch-extra`, which adds support for other architectures. This provides the program `qemu-system-aarch64`, which spins up an emulator for Aarch64-based machines. The RPi 3 is supported as a machine preset, using the command line argument `-machine raspi3`.

The emulator can emulate everything about the RPi except for the first stage bootloader, which is the code that runs at the very start of the boot process on a real RPi and then loads the kernel (`kernel8.img` for 64 bit mode). This bootloader is proprietary and runs on the RPi's VideoCore GPU. I am not trying to develop a first stage bootloader, though, so I do not need to worry about the lack of this functionality. The compiled kernel is passed as an argument to the emulator when we start it (`-kernel kernel8.img`). In the same way, we can give the SD card image to the emulator, with the argument `-drive file=sd.img,if=sd,format=raw`.

5.3.3 Debugging

GDB [11] is a debugger which can be used to debug executable programs. Similar to QEMU only supporting one architecture by default, GDB only supports running a program directly on the current machine. However, it does support remote debugging with a commands once it has started (`target remote URL:PORT`). Like with QEMU, there is an additional package which provides support for additional architectures (`gdb-multiarch`). When we launch QEMU, we can tell it to wait for a remote debugging session to connect before it starts the emulator (using the commandline arguments `-s -S`). Once QEMU is waiting for the debugger, we launch GDB (`gdb-multiarch kernel8.elf` in a terminal) and then connect to the emulator as a remote debugging client (using the GDB command `target remote localhost:1234`).

Now we can debug the OS kernel as if it is a regular program, using breakpoints and disassembling it as we run through the execution. This is how I debug the kernel when a new feature is not running as expected, and it allows me to determine the cause of any issues more easily than a debug shell would have done. Also, the debugger works without needing a working kernel, so I can debug issues which cause the kernel to fail.

5.3.4 Organising the build process

When I first started this project, I was using CMake to automate the build process. CMake is a tool which tries to automatically find the correct dependencies required for compiling a project (the compiler, linker, and also any libraries and included files can all be handled by CMake). Unfortunately, CMake was causing issues with compiler and linker paths, and it was not consistently finding the correct linker for the compilation of this project.

I then rewrote the build system using GNU Make (with a Makefile). This allowed me to hardcode the correct compiler and linker for the build process, and I haven't encountered any more issues with the compilation toolchain.

The top-level Makefile does not contain any rules to invoke the compiler directly, and only handles the linking together of the compiled object files into the final kernel. A separate Makefile is used to compile each source file, and Make invokes itself recursively in each source directory for the compilation.

The source files are kept separate from the object files and compiled kernel, by putting all the results of compilation into a `build/` directory at the top of the project (source files are located in `src/`). The object files and final compiled binaries are also separated, with object files being in `build/obj/` and binaries in `build/bin/`. Keeping the directory structure organised in this way helps maintain an understanding of the overall project structure as the project grows in size, as not keeping the order in this way would result in quickly becoming overwhelmed by the sheer number of files.

6 Design

6.1 Operating System (OS) components block diagram

Figure 2 is a diagram of how the different components of the Unix OS fit together. It shows how user programs running on the OS are supported by the OS components. The components mentioned include user libraries, as well as various parts of the kernel, such as system calls (see subsubsection 8.2.2), the Filesystem (FS) (subsubsection 8.2.4), and the *process control subsystem.*

Figure 2 is also a representation of one possible final goal state of the OS, specifically a final goal state that is very heavily influenced by Unix, but there will be different designs which the OS could follow moving on after this project's timescale. A diagram in the same style which more accurately represents the components that have been implemented during this project is shown in Figure 3.

The differences between the two designs, as well as one proposed path from this project's result to the *Unix structure*, can be seen in Figure 4.
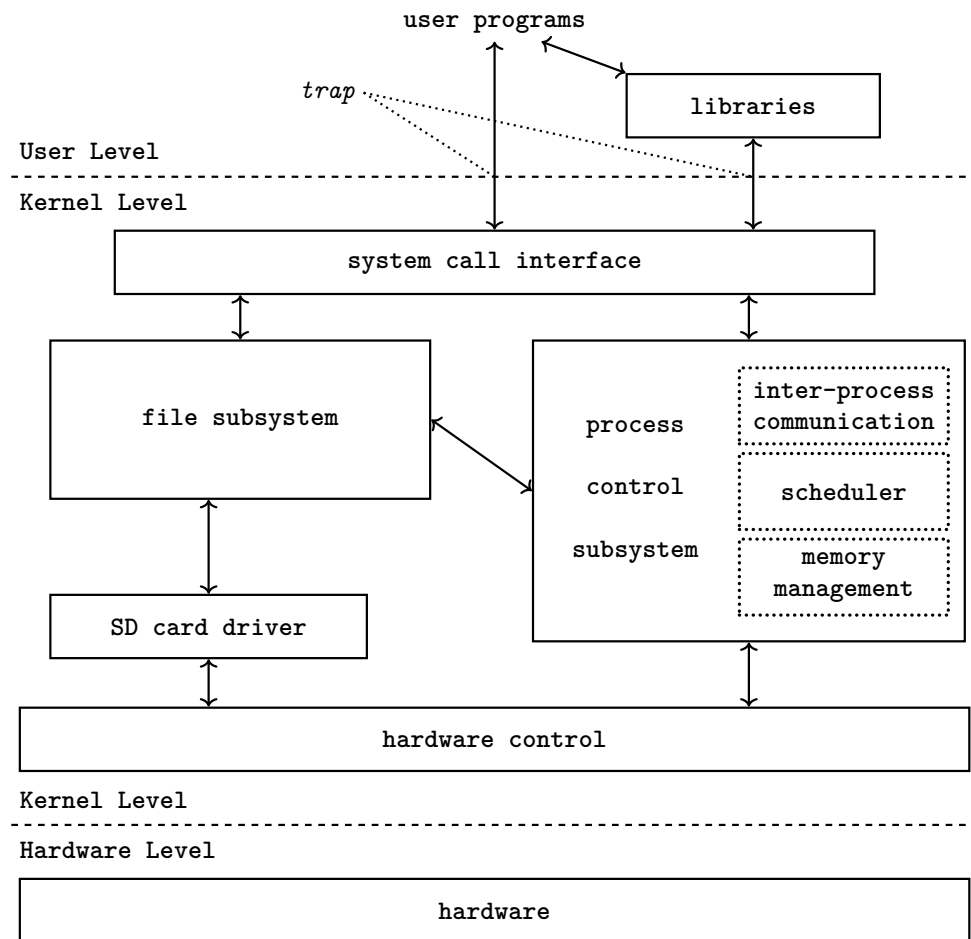
Figure 2: A diagram of how the components of the Unix operating system interact with each other. Inspired by a diagram from a book on Unix (Figure 2.1 from [3]). The versions of this diagram which are available online are low-resolution, so the version provided here is in a vector format, and this document's source code is available online [20].

6.2 User Interface and the graphics of the OS

The interface of the OS is the console. It is the text that is displayed to the user from the Raspberry Pi (RPi). I will implement some quality of life features, such as terminal scrollback, which allows the user to scroll up and see what was printed to the terminal earlier. This feature reduces the need for a terminal output pager such as `less`, which will not be implemented until very late in the development.

```
┌─────────────────────────────────────────────────────────────┐
│                        kernel shell                          │
└─────────────────────────────────────────────────────────────┘

          ┌─────────────────────────────────────┐
          │            file subsystem            │
          └─────────────────────────────────────┘

              ┌─────────────────────────────┐
              │        SD card driver        │
              └─────────────────────────────┘

┌─────────────────────────────────────────────────────────────┐
│                      hardware control                        │
└─────────────────────────────────────────────────────────────┘
```

Kernel Level
- -
Hardware Level

```
┌─────────────────────────────────────────────────────────────┐
│                          hardware                            │
└─────────────────────────────────────────────────────────────┘
```

Figure 3: The structure of this project's OS at the end of the project.

Figure 4: Three diagrams showing a suggested development path the OS could take after this project. The end goal of this path is the structure of the Unix OS. The three diagrams are shown full-size in the appendix.

11

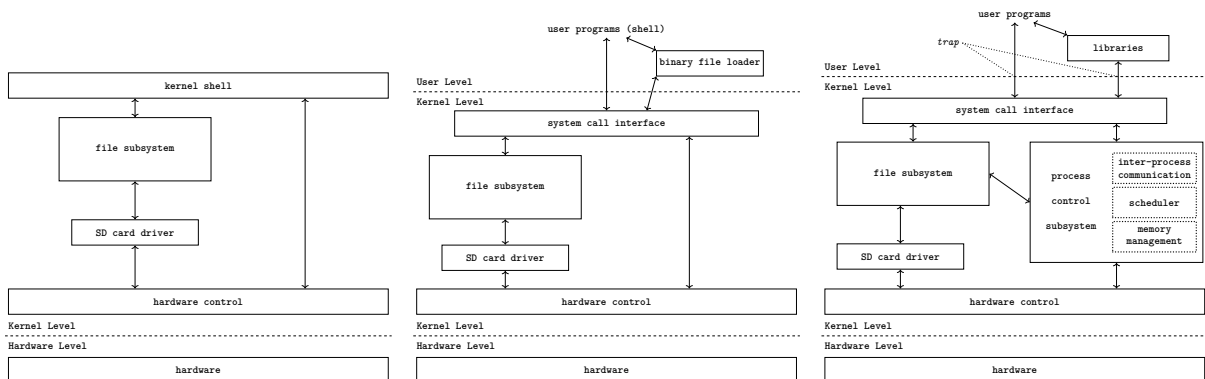The font I decided to use for the console is Bizcat [4]. It is an 8x16 bitmap font, and I converted it into a C source file as a constant array. The font is easy to read at a screen resolution of 640x480, which gives 80 columns and 30 rows of text.

The original colorscheme of the console was solid white text on a completely black background, but I found that this was too harsh to look at for long periods of time. I changed both the background and foreground colours to what I found more pleasant to read text from, with a very light grey text on a very dark grey background. A comparison of the colorschemes can be seen in Figure 5.


Figure 5: The original colorscheme of the console (left) and the new colorscheme (right).

7 Programming languages and compilation toolchain

7.1 Programming languages used

From the beginning of this project, I had already decided to write the majority of the code for the Operating System (OS) in C [24]. C is primarily a systems implementation language (making it perfect for writing an OS), and I was already very familiar with it, having used it regularly for the past 5 years.

For some very low-level parts of the OS, I needed to specify the exact order in which I wanted arm instructions to be executed, so ARMv8 assembly was the natural choice, being the closest language to the machine code that gets executed without actually being the machine code. Assembly language allows preprocessor macros, so I did not have to manually align instructions to 4-byte boundaries anytime there was a change, and header file "include"s which allow register values to be abstracted away from the logic and replaced with meaningful names. These features were used extensively in the bootloader and in subsubsection 8.2.3 for explaining the values placed into various Central Processing Unit (CPU) registers.

7.2 Problems with the original compilation toolchain

At the start of development, I was using the GNU cross-compiler for bare-metal arm 64-bit targets, `gcc-aarch64-none-elf`. This was working but the build system I was using at the time (CMake) was easier to configure using clang/llvm as a cross-compiler. Also, the GNU cross-compilers are not available in most Linux distributions' core repositories, so clang/llvm is much easier to install, needing only the package manager of a Linux system.

Later in development, I was having difficulties maintaining the build system which I had started the project with (CMake), as the intended usage is to automatically detect the executables that should be used for compilation (the cross-compiler, linker, etc.). This caused problems because if I modified the status of my build computer, by installing some new packages for example, CMake would detect and start using the wrong linker, causing the build process to fail*.

As a result of the problems I was having with CMake, I decided to use GNU Make (with a `Makefile` instead of a `CMakeLists.txt`, calling the `make` executable) instead. I created a Makefile for the project, and this allowed me to hardcode the name of the compiler executable, and all the tools needed to generate the

---

*I later discovered the cause of this. I had re-installed the `gcc-aarch64-none-elf` cross-compiler, which resulted in the GNU arm 64-bit linker appearing before llvm's `ld.lld`. CMake then tried to use `gcc` (**not** the cross-compiler) to invoke the linker, but the arguments for invoking the linker via `gcc` and `clang` are different, so `gcc` was giving back an error, saying that the `-T` argument was missing a parameter, and I was unable to force CMake to call `ld.lld` directly.

compiled kernel image and SD card image. During the implementation, as the number of source code files grew, I decided to start using sub-make [16] for compilation so that the main Makefile could be kept simpler, so I moved the compilation commands into a separate Makefile which handles each subdirectory of the main `src/` directory.

8 The Kernel

8.1 Design

8.1.1 Bootloader

On the Raspberry Pi (RPi) 3, the Graphics Processing Unit (GPU) contains a first-stage bootloader which automatically does most of the required setup for booting into an Operating System (OS) kernel. However, the Central Processing Unit (CPU) will not be setup completely, and so I will need to write some code to set register values and verify that everything is as it should be before we jump into the kernel `main` function. I will call this code the 'second-stage' bootloader.

8.1.2 UART driver for keyboard input and initial console output

Universal Asynchronous Receiver-Transmitter (UART) is used on the RPi for serial communication, with two different UART ports: a mini-UART and a full PL011 UART [15]. The PL011 UART is faster and more reliable than the mini-UART, but it is disabled by default, enabling it disables bluetooth, and it is more complex to program for. I will use the mini-UART to debug the setup of the PL011 UART, which will be used as the main serial connection once it is enabled and working. I won't be using bluetooth in this project, but future projects which expand on my code could switch back to the mini-UART if they wanted to use the RPi's bluetooth chip.

Serial communication is used in this project as a means of interacting with the OS, as I will not be implementing support for USB peripherals. All keyboard input is expected to come through the serial port. I will also use the serial port for debugging as I develop the graphics driver (see subsubsection 8.1.3).

For the initial console output, I will implement a `printf` function for formatted printing of characters, integers, hexadecimal numbers, and strings. This will allow me to easily print the values of registers, so that I can make sure that all values are as I expect during the boot process. The `printf` function will be generic, so that I can use the same `printf` function for both of the UARTs and also eventually for the graphical console.

The OS will get keyboard input from the serial port, which will be provided by QEMU, the RPi emulator. QEMU uses VT100 escape sequences [23] for its serial interface, so normal keys on the keyboard (and shifted keys) are passed to the OS as the ASCII value of the characters they represent.

8.1.3 Graphics driver for displaying pixels on the screen

Once I have a working serial terminal, I will implement a graphics driver for interacting with the RPi's framebuffer. This will include a function for setting individual pixels, a function for setting the entire screen to a colour, and a function to display a bitmap to the screen at a given location. With these functions, I will write functions for displaying characters to the screen, with variables for the current position on the screen. I will convert a bitmap font [4] into a constant array of bitmap values, and the character printing function can use the character as an index into this array to get the bitmap to display to the screen.

The screen blanking function will use a `memset` function to set the entire framebuffer to a given value, as this would be much faster than setting each pixel individually. The bitmap function will use the pixel setting function internally, and will use the bitmap to decide which pixels should be set. Each of these functions is very simple on its own, and will be relatively easy to implement, but once put together they will create the complex behaviour required for displaying characters to the screen.

8.1.4 System calls

TODO: explain that syscalls are split into 2 different components: the mechanism to allow for them and the implementations of the syscall functions themselves.

System call mechanism using software interrupts The kernel will include a system call mechanism, which would eventually be the way for processes (at the user level) to escalate privileges temporarily to request that the kernel perform some action on their behalf. Processes and the user / kernel level split will not yet be implemented, but I will implement support for system calls to prepare for the time when they are needed.

The OS will use software interrupts to trigger a system call, with an exception handler to interpret which system call function needs to be called. On arm CPUs, a software interrupt is capable of changing the Exception Level (EL). In user-mode (once that is implemented) the CPU would be in EL 0, and kernel code will execute in EL 1. The software interrupt will place the CPU into EL 1, entering kernel-mode, and when the system call function returns, the exception handler will run an exception return instruction, placing the CPU back into the EL that it was in when the software interrupt occurred. This is how the system call mechanism will allow the OS to separate user space from kernel space while still allowing user-mode processes to perform operations that require kernel permissions.

Each system call will be assigned a number, and I will copy these from Linux at the start. Most Linux system calls will not ever be implemented, but having some level of equivalence between this system and Linux system calls will ensure that developer will feel confident that at least some of the knowledge is transferable.

The system call exception handler will use the system call number to determine which function to call. This will be done using a table in memory which maps system call numbers to function pointers. The exception handler will lookup the correct table entry and then call the function at the address it finds.

System call implementations The implementation of the system calls can be separated from the mechanism which enables them to operate. TODO: Expand this

8.1.5 SD card driver

In order to read data from the SD card on the RPi (and write data back to it), the OS will need to implement a driver for the RPi's eMMC chip. The simplest way to do this will be to copy an existing driver for that chip. I will take code from one of the related works [18].

8.1.6 exFAT filesystem

The kernel will include an implementation of the exFAT filesystem. I considered several different filesystems including FAT32, exFAT, and YAFFS [2]. The exFAT filesystem's popularity was the reason I chose it over YAFFS. The FAT32 filesystem seems simpler to implement than exFAT, but I was not able to find a specification for it online, whereas the exFAT official specification is available on Microsoft's documentation website [10]. Having a single source for the filesystem specification gave me more confidence that different tools would work together. Also, the exFAT filesystem is partially supported on Linux systems (FAT32 is not as well supported), and I am developing the OS on Linux.

The file system implementation will need several features in order to work as the main filesystem of this OS. These are listed below, using terminology that will be explained in subsubsection 8.2.4.

- Read the 'superblock' of the SD card

- Interpret the FAT table to follow cluster chains

- Read directory entries

- Search a directory for a file according to its name

- Read the contents of a file

- Write to a file and correctly update its metadata for size and modification time

TODO:
Nice-to-haves

- Virtual filesystem abstraction

- System call implementations for commonly used operations

- Kernel / user mode separation

Future possibilities - move to evaluation section

- Scheduler for multiprocessing (with a process control table / vector in the kernel)

- USB driver for keyboard input (a port of an existing USB driver project)

- Drivers for some other hardware on the RPi (Ethernet / Bluetooth)

8.2 Implementation

8.2.1 Bootloader

The second-stage bootloader* is written in arm assembly code. This is because the bootloader performs the necessary setup in order to get the CPU and memory into a state in which code written in C can be run. Once the bootloader has finished this setup, it jumps into the kernel main function (written in C) straight away.

First, the bootloader stops all but one of the CPU cores. The current state of the OS does not support simultaneous multiprocessing with the multicore CPU on the RPi. This first step effectively converts the 4-core CPU into a single-core CPU by putting cores 1, 2, and 3 into infinite 'do-nothing' loops, leaving only core 0 to execute code.

Next, the boot code gets the current EL from the CPU. The EL is stored in a special CPU register on arm CPUs, called `CurrentEL`. On the RPi this should always be 2, but I also wrote code to handle the case where it is 3. If the EL is 3 then the boot code sets the needed registers and safely transitions the CPU into EL 2 in the same way as described below.

Changing the Exception Level Once the boot code verifies that it is executing in EL 2, it prepares to change to EL 1, the level that the kernel code will run in. It sets important registers, such as the exception vector table (used by the system call mechanism) and the exception return address (the address of the code to execute once the CPU switches to EL 1). It also ensures that EL 1 is executing in 64-bit mode.

Once all the parameters are set correctly, the bootloader executes an exception return instruction (`eret`). Exceptions are the only way to increase the EL on an arm CPU, and returning from exceptions is the only way to reduce the EL. The exception return instruction changes the CPU into EL 1 and resumes execution at the address of the EL 1 code.

The floating point / vector registers Once I started to implement the function `printf`, which is a *variadic* function†, I started to experience strange behaviour. I debugged the compiled kernel using remote debugging in GDB, and discovered that upon entering the `printf` function, the executing kernel would jump to an address near `0x0`, the beginning of memory. This meant that an exception was occurring and the CPU was going to a position in the *Exception Vector Table*. I determined from the offset into the exception table that the exception was permission related.

Looking at the first instructions in the disassembly of the `printf` function, I saw that the function was accessing registers I did not recognise, labeled `Qn` instead of `Xn`. After some research I determined that the `printf` function was trying to use the 32 floating point registers of the CPU to store the variadic argument list, but the floating point registers were disabled by default.

The solution was to enable the floating point registers after switching to EL 1, by setting some bits in one of the system registers. Setting those bits enabled the floating point registers (also called vector registers) and enabled the use of variadic functions in the C code.

8.2.2 System Calls

The system call interface required writing an exception handler and an *Exception Vector Table* The *Exception Vector Table* is a table in memory which gives the CPU instructions to execute for each different type of exception that can occur while it is running. These could include hardware interrupts such as timers, software interrupts (like the system calls we are interested in), or more serious errors like invalid instructions or insufficient permissions for an operation. I originally wrote the Exception Vector Table using outdated documentation, so the padding between entries was incorrect. I fixed this issue, and placed a jump instruction in the software interrupt exception entry. This branched the CPU to the exception handler that I wrote.

I wrote the exception handler for software interrupts in arm assembly code, the same as the Exception Vector Table. First, it checks that the instruction that generated the exception was the supervisor call instruction, `svc`. Then it checks the argument of that instruction – we only want exceptions where the argument was 0. After checking that the exception was indeed a system call, the exception handler gets the system call number and the (up to) 6 arguments and places them into the first 7 general purpose registers. It reads an index into a named table (`sys_call_table`) and uses the address it reads as a function pointer, calling the function. This is the mechanism used to identify the correct system call function from just the system call number.

Also in the assembly code for system calls is an indirect system call function (`syscall`). This function is equivalent to the Linux function `int syscall(long number, ...)`, and is the intended way for C code to

---

*The 'bootloader' is more accurately described as boot code, since it does not do any loading of the kernel from disk, the traditional task of a bootloader.

†A variadic function is one with a variable number of arguments

```
extern int syscall(long nr, ...);
// read syscall is syscall number 0
// fd is a file descriptor
// buf is a void pointer to a buffer large enough to read to
// count is the number of bytes to be read
syscall(0, fd, buf, count);
```

Listing 1: How some C code would invoke the system call `read` in the current implementation.

```
struct exfat_superblock {
    uint64_t partition_offset;
    uint64_t volume_length;
    uint64_t sectorsize;  // Bytes per sector
    uint64_t clustersize;  // Sectors per cluster
    uint32_t fat_offset;
    uint32_t fat_length;
    uint32_t clusterheap_offset;
    uint32_t cluster_count;
    uint32_t rootdir_start; // First cluster of the root directory
    uint32_t volume_serialnumber;
    uint16_t fs_revision;
    uint8_t  fat_cnt;  // Number of FAT tables
    uint8_t  use_percent;
    // Volume flags
    uint8_t active_fat;  // Which fat is active
    bool volume_dirty;
    bool media_failure;
    bool clear_to_zero;
};
int exfat_read_boot_block(struct block_device *, struct exfat_superblock *);
```

Listing 2: The structure used to represent the exFAT superblock information, and the declaration of the function used to read it

invoke a system call. Eventually I would write named wrapper functions for the common system calls, but for now the way that C code would run the `read` system call, for example, would be as shown in Listing 1.

8.2.3 CPU setup

When the kernel starts, the CPU will either be in EL 2 or EL 3 (The RPi only supports EL 2 but I have written code to support EL 3 just for completeness). The boot code of the kernel checks the current EL and runs some code to put the CPU into EL 1. The code sets up the system registers correctly to fake the CPU returning from an exception. When the registers are correctly set, the arm instruction `eret` is run, and the CPU "returns" from the "exception handler", jumping into the kernel code and entering into EL 1. This is done in this way because exceptions are the only way to change EL, and the only way to reduce the EL is to perform an exception return (`eret`).

8.2.4 exFAT Filesystem (FS)

I took an existing driver for interfacing with an SD card (via the RPi's eMMC controller) and integrated it into the kernel. This code defines a structure called a `struct block_device`. I also copied some FS code from the same related work [18], but most of the code for the FS I wrote by myself.

I first created a function to read the 'superblock' of an exFAT system. The superblock is the first block of the disk partition, and contains some critical information for interpreting the information on the rest of the disk. The structure used to store this information and the declaration for the function are both sown in Listing 2. The function is called `exfat_read_boot_block()` because sometimes the superblock is referred to as the 'boot block' (although ususally only when talking about disks instead of FSes).

There is also a function which takes a cluster index and returns the index of the next cluster in that cluster's cluster chain. This information is retrieved by reading the File Allocation Table (FAT) section on the volume (see section 4.1 in [10]). The function signature is shown in Listing 3.

```
static uint32_t fat_idx_to_cluster_idx(struct block_device *dev, struct
         exfat_superblock sb, uint32_t fat_idx);
```

Listing 3: The function which is used to get the cluster index of the next cluster in a cluster chain

```
struct exfat_directory_info {
    struct block_device *bd;
    struct exfat_superblock super;
    uint32_t start_cluster; // First cluster of the directory
};

struct exfat_dirent_info {
    struct exfat_directory_info *parent;
    uint32_t direntset_idx;  // Index into the "Directory"
    // The importance, category, and type code represented in the same way as
    // on disk.
    enum EXFAT_DIRENT_TYPE type;
    bool in_use;  // The in_use bit of the entrytype field can change
    union {
        struct exfat_dirent_bitmap_chunk bitmap;
        struct exfat_dirent_upcase_chunk upcase_table;
        struct exfat_dirent_vollab_chunk volume_label;
        struct exfat_dirent_file_chunk file_metadata;
        struct exfat_dirent_volguid_chunk volume_guid;
        struct exfat_dirent_strext_chunk stream_extension;
        struct exfat_dirent_fname_chunk filename;
        struct exfat_dirent_vendor_guid_chunk vendor_guid;
        struct exfat_dirent_vendor_alloc_chunk vendor_alloc;
    } entry_data;
    // If the dirent doesn't point to another cluster chain, then data_length
    // will be 0.
    uint32_t first_cluster;
    uint64_t data_length;
};
```

Listing 4: The structures which describe directories and directory entries on an exFAT volume

The next structure on an exFAT volume that needs to be interpreted is the directory. Director-ies primarily contain information about files, but can also contain some special entries. These are all stored in 'directory entries' (see sections 6 and 7 in [10]). The information for a directory is stored in a `struct exfat_directory_info`, and information describing a single directory entry is stored in a `struct exfat_dirent_info`. These are both defined in Listing 4.

A function in the exFAT driver, `exfat_read_directory_entry()`, is used to read and interpret a single directory entry. Its single argument is a pointer to a `struct exfat_dirent_info`, which is already filled with the information that is needed to identify a single directory entry in a directory. The function fills the rest of the fields.

Another function, `exfat_readdir_fromblock()`, takes a block device (the SD card) and the address of a block which is the beginning of a directory, and reads the entire directory. It stores each file from the directory in a `struct dirent`, partially defined in Listing 5. The `dirent` structure contains a `next` pointer, so the function joins all of the files into a linked list data structure, and returns the first file.

A distinction is made between a directory entry and a file. A directory entry may not be associated with a file, and each file will have at least three directory entries associated with it. The `exfat_dirent_info` structure refers to a single directory entry, whereas the `dirent` structure describes a single file. The `exfat_readdir_fromblock()` function reads directory entries in a loop, ignoring those not associated with files. It reads all of the entries for each file, then creates the `struct dirent` and appends it to the linked list. When it reaches the 'end of directory' directory entry the function returns.

```
// Struct representing a file
struct dirent {
    struct dirent *next;  // The next file in the directory
    char *name;  // Filename
    uint32_t byte_size;  // Size of the file
    uint8_t is_dir;  // Is this file a directory?
    ...
};
```

Listing 5: The dirent struct, which stores information about files

8.2.5 Graphics driver

For graphics, I created a simple library for manipulation of the framebuffer. In order to print characters to the screen, I converted an 8x16 bitmap font [4] into a constant array definition in a C source file. I then wrote a function to print a character from this array onto the framebuffer, and I added global variables which keep track of the current character position on the screen.

TODO: write about changes here (scrolling at the bottom of the screen - uses the swap rectangle function)

When the console reaches the end of a line, it moves to the start of the next one. At the moment, when it reaches the end of the last line, it moves back to the start of the first line. In the future I will implement scrolling, so that when the console reaches the bottom of the screen, every line will be moved up to make space for the next line.

8.3 Evaluation

Overall, the kernel met the requirements that I set out in subsection 8.1. Most of the components designed for the kernel have been implemented. The ones which have not been fully implemented are either almost finished and lacking some other piece of the OS, or the groundwork has already been laid to make developing the features easier. This was done to help new developers to more easily understand the source code of the project.

8.3.1 Bootloader

The bootloader is robust and reliable, starting the kernel main function every time the OS boots. It sets up the CPU into the correct configuration, following the kernel's current requirements.

8.3.2 System call mechanism

The interface to system calls is well-implemented. The wrapper function used to invoke a system call checks that the system call number is within the valid range of numbers, preventing unnecessary software interrupts from being generated. The complex mechanism required for properly managing system call generation and handling has been broken down into logical parts which are each small enough to be easily understood.

8.3.3 System call implementations

There are currently no system call functions in the OS. The table which is used by the system call exception handler is in place, so once the system call functions are written they can easily be dropped into the table and then immediately those system call numbers will start to function.

TODO: refer back to second aim of the project.

8.3.4 Filesystem

The exFAT FS was only just implemented. Reading directory entries works, but there aren't many functions built on top of that to allow for more abstract access to the FS. The virtual filesystem abstraction is severely lacking in the current version of the OS. All code which interacts with the FS directly uses exFAT-specific functions. This limits both ease of programming with the FS, and the overall usefulness of the FS.

I started to move towards having a simple virtual FS abstraction in recent development by creating several of the functions presented by the `struct fs` structure. These include a version of `fopen`, `fread`, and `fclose`, which are functions needed for reading files.

I started to write the exFAT-specific version of `fread` very recently in development. The function declaration is shown in Listing 6, along with versions of `fopen` and `fclose`. The arguments of the functions

```
FILE *exfat_fopen(struct fs *, struct dirent *, const char *mode);
size_t exfat_fread(struct fs *, void *ptr, size_t byte_size, FILE *stream);
int exfat_fclose(struct fs *, FILE *fp);
```

Listing 6: Function declarations of the exFAT versions of `fopen`, `fread`, and `fclose`.

are different from the standard arguments to these functions. This was changed to make the implementation of the functions easier, because the `struct dirent` already points directly at a file, whereas a string representing a file path can be difficult to interpret. Unfortunately, this change just made it more difficult to write code which uses the functions, since the calling function now needs to do the job of converting a path into a `dirent` structure.

8.3.5 Testing

The test coverage of the kernel is currently very limited. I wrote enough tests to determine that the testing system works and intended, then after that no more tests were written. I should have been writing unit tests for every kernel component, but instead most of them are effectively untested.

The lack of testing provides a great opportunity for new developers to become familiar with the project. Writing tests for an existing component would give a new developer very good in-depth knowledge of how that component works. A developer who enjoys writing unit tests would be able to learn the entire kernel just from writing tests. The testing could even be expanded to components of the shell, which would allow a new developer to both improve the reliability of the OS and learn how it works, priming them to start creating new features.

8.3.6 Kernel / user space

The OS features no separation between *kernel space* and *user space*. Everything that runs under the OS is effectively kernel code, including the shell. Ideally, the shell would run at a lower EL than the kernel (i.e. EL 0), giving it reduced permissions to access hardware and execute restricted instructions (such as stopping the CPU). However, doing this would require system calls to be available to the user mode shell, and (as mentioned in subsubsection 8.3.3) there are currently no system call functions implemented in a working state. Therefore, creating a separation between kernel code and user code is not currently possible. This will be a high priority feature once system calls are available.

On the other hand, this feature could be used to drive the implementation of system call functions. A developer could start work on getting non-critical code to run in EL 0, and then implement system calls as they are required. This would be possible because the system call mechanism is already implemented, and the only thing needed to get a given system call working is an implementation of the function. A small amount of experimentation will also be required to discover if any additional changes need to be made to the boot code to get the CPU to run in EL 0 correctly.

9 The Shell

TODO: reword this into an introduction to the shell

- Design of the shell is important

- If it is done well people will immediately feel comfortable using the Operating System (OS)

- A simple but intuitive shell will give a good first impression to the project's code

FIXME: The shell which I will include in the OS will be a port of a POSIX compliant, simple shell, such as Dash [8]. This shell will be the main way that users interact with the system, as it will allow them to execute programs which do anything they need, such as compiling other programs, editing text files, or viewing the contents of files. The shell will run in user space, and any time it needs to run kernel-level functions, it will do so through the system call interface. This is the same way the Bourne Shell ran on Unix, and it is a known good way of allowing the user to control a computer. Anyone who has experience with a shell will immediately feel at home when they see the input prompt and a flashing cursor.

9.1 Design

9.1.1 Commands as functions

In the initial implementation, I will develop the shell as a part of the kernel, and all of the commands will be functions built into the final compiled kernel. I will do it this way because it is simple, and my OS will not have a binary file loader at this stage, which would be needed for loading external versions of commands. The downside of using this method is that any new commands or changes to commands would require a recompile and rebuild of the entire kernel, but so far building the kernel from scratch is a relatively quick process.

Commands in the shell will be functions in the shell source code files. The functions for commands will all share the same function signature, which will be the same as the function signature for a `main` function [13]. This will mean minimal changes would be needed to eventually turn each command into its own executable, with the binaries being stored on the filesystem (in a `/bin` directory). This change will be the natural next step from what I create, so I am designing my shell implementation with that in mind from the start.

9.1.2 The shell loop

The logic of an interactive shell is quite simple. It consists of an infinite loop:

- Read a line from the keyboard

- Split the line into words by spaces

- Determine which command, if any, the first word matches

- Execute the correct command

- Repeat this process forever

The shell needs to get user input on a line-by-line basis, since each command starts at the beginning of the line and ends at the newline character. In the first iteration, I will store the text input in a fixed-size buffer, meaning that there will be a maximum size for an input line. I will set this maximum size to a reasonable value such that it will not cause any problems, and I will leave this problem to fixed in a future update to the shell's user input functions.

I will split the line by spaces and place the words into a 'vector', passed to the commands as the argument `argv`.

9.1.3 The commands

I will create several commands, starting with the simplest and most essential ones. The OS will need a way to shutdown the Raspberry Pi (RPi), so I will implement a `shutdown` command. I will write a `pwd` command to print the current working directory, as this will be a very simple (and also useful) operation. When I write the filesystem, I will create a function for printing the contents of directories (`ls`), and also a recursive version of this, `tree`, shich shows an entire directory structure in context.

Once I have implemented enough of the filesystem to get file contents from the SD card, I will create a `cat` command to print files to the screen. Finally, a command for getting file attributes, like `stat` on Linux systems. On Linux systems, `stat` is implemented via a system call (`fstatat`), so this command could be rewritten using that if a future development creates that system call.

9.1.4 Executable file loader

One feature that would significantly improve the shell would be the inclusion of a binary file loader. This could either support only raw binary files, or statically compiled Executable and Linkable Format (ELF) files. A file loader would not be possible without first implementing file reading in the filesystem, so this would be a very late feature to start implementing if I put it into the OS.

For loading ELF files, I would implement a basic ELF parser to extract the instructions from the file and put them into memory. The loader would then determine where the entry point is, setup the stack and Central Processing Unit (CPU) registers, then finally jump to the entry point.

The process would be the same for loading raw binary files, except that the entry point would be fixed. There would also be no need to extract the instructions from the ELF structure, just loading the instructions from the filesystem and then jumping into the loaded file contents.

```
struct shell_command {
    int (*fun)(struct command_args);
    const char *cmd_str;
};
```

Listing 7: The definition of the `shell_command` struct

9.1.5 Rewrite commands using system calls

Many of the commands I will implement would benefit from using some standard system calls, but I will not
have implemented any system calls by the time I create the commands. Rewriting these commands to use
the system calls would be a nice-to-have feature. It would also give an incentive to implement the common
system calls, as I would be developing system calls I need at the time.

Some commands which would benefit from this would be `ls` (system call `getdents`), `stat` (`fstatat`),
and `cat` (which would benefit from the system calls `open` and `read`). The `printf` function could also be
rewritten to use the system call `write`, writing to a file descriptor which represents the graphical console.

9.1.6 Replace the shell and all commands with separate executables

As a follow up to including a binary file loader in the shell, I could also completely separate the shell and
kernel. This would allow me to modify the shell, which will be growing in size, without recompiling the
kernel, which will also be growing in size. The individual compile-times of both the shell and the kernel
would be effectively cut in half, allowing for much more rapid iteration and development of both.

With the shell separated from the kernel, the kernel would load the shell from the filesystem on boot
(using the binary file loader). The commands could also be moved to their own executables, and placed into
some common directory, likely `/bin/`. The shell would search this directory for filenames when it is trying
to find executables that match input lines. This hardcoded directory would eventually be replaced by an
environment variable `PATH`.

Doing this would allow modification of the commands available on the system without recompiling the
kernel, restarting the OS, or even restarting the shell.

9.2 Implementation

9.2.1 Getting input from the keyboard

The shell needs to get input from the keyboard line-by-line, so the terminal can buffer the current line as
it is being typed. A buffer of a fixed size is created on the stack when the shell is started. The terminal
uses a `getc()` function to get one character from the serial port. When `getc()` returns a newline character,
the shell's input line buffer gets tokenised, splitting it by spaces. The function `strtok()` (from `string.h`)
is used for this.

The addresses of the strings containing the tokenised words from the line are placed into an array, which
is the member 'argv' of a `struct command_args`. The total number of words in the line is placed into that
struct's member 'argc'. The member variable 'envp' is set to a pointer to the global `struct process_env`
of the kernel. This struct contains all the miscellaneous information that commands may need, such as the
current working directory and the return value of the previous command. Once the OS implements the
concept of user and kernel space, only a subset of this struct would be passed to commands, and then the
commands would need to use system calls (see subsubsection 8.2.2) to access that information.

9.2.2 Storage of enabled commands

The shell's memory contains a linked list which contains all of the enabled commands. The commands are
described by a structure called a `struct shell_command`, with the structure definition shown in Listing 7.
The member `cmd_str` is a string containing the word that is used to invoke that command, and `fun` is a
function pointer to the function that will be used whenever the first word of an input line is equal to `cmd_str`.

9.2.3 Shell commands as functions

In the current implementation, all of the commands that are available to the shell are functions in the
source code of the kernel. The function signature of the command functions is shown in Listing 8, as
`command_function()`. The function `add_command()`, also declared in Listing 8, takes a string and a function
pointer and creates a `struct shell_command` from them, adding it to the shell's `struct shell_cmdlist`

```
int command_function(struct command_args);
int add_command(const char *cmd_str, int (*fun)(struct command_args));
int shell_init();
```

Listing 8: The signatures of some important functions for the shell

(the linked list). The function `shell_init()` calls `add_command()` once for every shell command, initialising the shell to prepare it for running commands.

9.3 Evaluation

Overall I am pleased with the progress I have made with the shell, but I have not made as much progress as I originally hoped for. I think this was due to the kernel containing some surprises that I did not consider or have the knowledge to pre-empt at the beginning of the project.

The modular design for creating shell commands was a good implementation choice. It is very simple to create new commands and get them working with the shell – this will be enticing to developers who see how easy it is to create a simple utility function[*]. All that is needed is a wrapper function with the correct signature and a name for the command.

The argument to the shell command functions is very similar to the arguments that a `main()` function would take once the shell moves to the commands being individual executables. The first two arguments[†] are exactly the same (`int argc, char *argv`). This similarity will make the transition to commands as separate executable files much easier since the command function can be copied directly to a main function and then compiled on its own.

The third member of the `command_args` structure, `envp` is badly named – it doesn't contain any environment variables. In the future, the data which is currently stored in this `struct process_env` will be stored in a table in the kernel, and tracked on a per-process level. Shell commands will use system calls to access the same information.

One command that I would like to implement in the near future is a basic text editor. This would be a line-oriented editor, with the basic functionality of `ed` [9], the Unix text editor. Writing a simple text editor will require the ability to write and create files, so this would prompt me to implement that functionality into the exFAT Filesystem (FS).

A final extension to the shell would be to include some shell language features. This will effectively make the shell program an interpreter for an `sh`-like language. The interactive version of the shell would become a Read-Evaluate-Print-Loop (REPL) for the interpreter. Implementing a full interpreter is obviously well outside of the scope of this project, but adding these features would turn the shell into something that could be used for scripting, which would complement the OS nicely, and would make it seem like a much more complete package. The inclusion of a scriptable shell would also be an attractive feature to many would-be new kernel developers, which links back to one of the main aims of this project.

10 Evaluation and External Aspects

10.1 Code style consistency

The consistency of my coding style was not good over the course of this project.

At the beginning of the project, I was writing functions with return type `ERROR_TYPE`, which was an alias of `size_t`. Every possible error was assigned a unique value (in an unmaintainable way), and for each component there would be a unique success value.

Once `errno.h` was created, I started using `errno` to represent errors. Functions would not return `int`, with a zero meaning success and a nonzero return value indicating a problem.

These two coding styles have co-existed in the project, but I should have refactored all of the `ERROR_TYPE` functions to return `int` and use `errno`.

10.2 Functions from `string.h`

Most of the functions declared in the `string.h` header file have no implementation. This is not a bad thing, as I was implementing those functions as I needed them. Most of the functions from that header file are very simple and take less than half an hour to write.

---

[*]A *Hello world* example function could be placed into the `README.md` of the git repository to showcase how easy it is.
[†]The first two members of the structure

Eventually I will reach a point where a very small number of the functions are missing implementations, and at that point I will either implement the rest of them all at once, or remove them.

The lack of simple function implementations could also be seen as a positive for the project. Providing implementations for these simple functions could be a non-threatening introduction to the project for an inexperienced developer, who may then go on to learn Operating System (OS) development using this project.

11   Summary and Reflections

The result of the project was not what I originally had hoped to achieve, but the development of the Operating System (OS) provided more education than I expected about the many components required for an OS to function. I believe that this educational experience will prove valuable to others, and so I consider this project to have been a success.

11.1   Project management

11.1.1   Re-planning the project

The original plan for the project was a little over-ambitious, which I discovered as I was developing the current implementation. I modified the plan considering my experience developing the OS so far. Some of the later, more complex tasks were removed from the plan, as I realised I would not be able to complete them, and some tasks were given more time.

11.1.2   Unexpected difficulties in development

I encountered several issues during development which I did not account for in the original project plan. The first is not so much an issue as it is a task not being required in the way that I expected. The **debug shell** was not needed in the way I had planned to implement it. My plan was to prompt the user for a command, and for the shell to support a very limited range of commands, perhaps starting with a command to print the contents of a memory address. This type of debugging has not been needed so far during the development process, and I don't believe it will be needed at all, as I can use The GNU Project Debugger (GDB) for debugging memory locations, and GDB doesn't even require the OS kernel to be in a working state, which would be the most likely scenario for if I would need to read memory.

The "debug" part of the debug shell task was still completed, as the current version of the kernel prints debugging information to the serial port (and also to the Raspberry Pi (RPi) video output once the framebuffer is allocated).

System calls were far more difficult to implement than I expected. The way that system calls were implemented was using arm exceptions to execute a function with the processor in a privileged mode of operation. Unfortunately, when I started working on the system call implementation, I did not have a very good understanding of the Central Processing Unit (CPU) mechanisms by which arm processors handle exceptions. The CPU uses a large number of system registers, along with the current Exception Level (EL), to decide where the exception handler can be found, and in which EL the CPU should handle the exception.

I spent some extra time learning about arm system registers in depth, and gained the knowledge I needed to finish the system call mechanism.

When I started working on the **filesystem**, I immediately realised something that I forgot to consider for the original project plan. The filesystem would require an SD card driver in order to interact with the RPi storage. I used an existing implementation of a simple RPi SD card driver [18], which only required a small amount of modification to get working. Once I had it integrated into my kernel, I created an SD card image and added it to the emulator, then tested the driver to see if it recognised the SD card.

11.2   Contributions and reflections

The project's external sponsor (who is also my project supervisor), introduced me to several of the related works [14, 17]. He also gave recommendations for some of the areas I was originally unsure of or less knowledgeable about. He suggested that I implement the FAT Filesystem (FS) as it is simple and relatively easy to implement, and he reminded me that on arm based systems, system calls would utilise software interrupts.

I think that the original project plan was too ambitious. I failed to consider that I would have other modules and also extra-curricular obligations, and I underestimated how much of my time would be taken up doing things unrelated to the project. Overall, the project has not progressed as quickly as I had originally expected. This is why I have scaled back the plan for the second half of the year, replacing the shell and libc with ports of already existing versions.

Acronyms

arm Arm Microprocessors Ltd.

CPU Central Processing Unit

EL Exception Level

ELF Executable and Linkable Format

FAT File Allocation Table

FS Filesystem

GDB The GNU Project Debugger

GPU Graphics Processing Unit

IO Input / Output

OS Operating System

RPi Raspberry Pi

SBC Single Board Computer

SOC System-On-a-Chip

UART Universal Asynchronous Receiver-Transmitter

References

[1] *A generic and open source machine emulator and virtualizer*. URL: https://www.qemu.org (visited on 2021-12-12).

[2] *A Robust Flash File System Since 2002*. URL: https://yaffs.net/ (visited on 2022-04-21).

[3] Maurice J. Bach. *The Design of the UNIX Operating System*. Prentice-Hall, Inc., 1986. ISBN: 978-0132017992.

[4] *Bizcat - an 8x16 bitmap font*. URL: https://robey.lag.net/2020/02/09/bizcat-bitmap-font.html (visited on 2021-12-12).

[5] *Clangd: teach your editor C++*. URL: https://clangd.llvm.org/ (visited on 2021-12-12).

[6] *ClangFormat*. URL: https://clang.llvm.org/docs/ClangFormat.html (visited on 2021-12-12).

[7] *CoC.nvim: the Conquer of Completion*. URL: https://github.com/neoclide/coc.nvim (visited on 2021-12-11).

[8] *Debian Almquist Shell (Dash)*. URL: https://git.kernel.org/pub/scm/utils/dash/dash.git/ (visited on 2021-10-26).

[9] *ed - edit text*. URL: https://pubs.opengroup.org/onlinepubs/9699919799/ (visited on 2022-04-21).

[10] *exFAT file system specification*. URL: https://docs.microsoft.com/en-us/windows/win32/fileio/exfat-specification (visited on 2022-04-21).

[11] *GDB, the GNU Project debugger*. URL: https://www.sourceware.org/gdb/ (visited on 2021-12-12).

[12] Brian W. Kernighan and Rob Pike. *The UNIX Programming Environment*. Prentice-Hall, Inc., 1984. ISBN: 978-0139376818.

[13] *Main function - cppreference.com*. URL: https://en.cppreference.com/mwiki/index.php?title=c/language/main_function&oldid=139269 (visited on 2022-04-20).

[14] *NetBSD*. URL: https://github.com/NetBSD/src (visited on 2021-11-28).

[15] *Primary and Secondary UART*. URL: https://www.raspberrypi.com/documentation/computers/configuration.html#primary-and-secondary-uart (visited on 2022-04-19).

[16] *Recursive use of make*. URL: https://www.gnu.org/software/make/manual/html_node/Recursion.html (visited on 2021-12-05).

[17] *RiscOS*. Version 5.29. URL: https://www.riscosopen.org/content/ (visited on 2021-11-28).

[18] *rpi-boot: A second stage bootloader for the Raspberry Pi (RPi)*. URL: https://github.com/jncronin/rpi-boot (visited on 2021-11-28).

[19] *The OSDev Wiki*. URL: wiki.osdev.org (visited on 2021-10-20).

[20] *The source code (in LATEX) of the document you are reading*. URL: https://github.com/Ytrewq13/YtrewqOS/tree/master/Documentation/dissertation/interim-report (visited on 2021-12-05).

[21] *Unix-like*. URL: https://en.wikipedia.org/w/index.php?title=Unix-like&oldid=1074888381 (visited on 2022-04-15).

[22] *Vim - the ubiquitous text editor*. URL: https://www.vim.org (visited on 2021-12-11).

[23] *VT100 escape codes*. URL: https://espterm.github.io/docs/VT100%20escape%20codes.html (visited on 2022-04-19).

[24] Brian W.Kernighan and Dennis M.Ritchie. *The C Programming Language*. Second. 1988. ISBN: 978-0-13-110362-7.

Appendices

A Interim: Appendices

A The original project plan

This appendix contains all of the tables, figures, and diagrams from the original project proposal, laying out my initial beliefs about what the project entailed, and how much work each part of the project would take.

The diagrams are copied directly from the project proposal, so tenses may now be incorrect (e.g. "I think" instead of the now more correct "I thought"), but this is done intentionally as this section is included only as a reference to the content of the previous document.

B Work tasks

Tasks are the individual components of the project, separated into the components of the Kernel and Shell, the miscellaneous parts, and the documents I will need to produce for the project. These tasks are enumerated in Table 3 and then assembled into a Gantt chart in Figure 6.

Task	Duration (weeks)
Kernel	
Bootloader	1
Graphics driver	2
Syscalls	6
Filesystem	3
Statically linked ELF Loader	2
Memory management	4
Process control	8
Init system and service manager	5
Threads and multithreading	4
Shell	
Debug shell	2
`pwd`, `cd`, `ls`, `stat`, etc.	2
`export`, variables, `set`, etc.	2
`if`, `while`, `for`, `case`, globs, etc.	3
Command substitution	4
IO pipes and output redirection	1
Misc parts	
Setup cross-compiler and build environment	1
IO and strings libraries	2
libc	8
`cat`, `head`, `less`, etc.	2
`roff`	1
`man`	1
Documents	
Interim report	1
Documentation	5
Dissertation	7

Table 3: The original plan of work for the project, including how many weeks I thought each component would take.

C   Gantt chart

Figure 6: A Gantt chart breaking down the proposed schedule of work for this project.

D   Work Packages

A "Work Package" is a collection of tasks with a well-defined deliverable. Not every task will be part of a work package, and some tasks will be work packages on their own. Table 4 contains all of the work packages for this project.

Work package	Description of the deliverable
Debug shell using the graphics driver	A debug shell which is printed on the Raspberry Pi (RPi)'s display and which uses keyboard input. The debug shell should accept several commands for debugging the processor status (including, but not limited to, printing the contents of memory, writing to memory, and jumping to a given memory address).
A basic kernel which can load (from disk) and jump into a compiled Executable and Linkable Format (ELF) program.	A program that can load a statically compiled ELF executable binary into memory and then jump to its entry point. This requires a filesystem library to be implemented, and also includes the ELF loader task.
Scheduling algorithm	A scheduling algorithm with support for multiple processes running on the system concurrently. This will include an implementation of `fork(3)` or a similar function in order to spawn new processes and `execve(2)` to replace the current program with a new process.
Init system and service manager	An init system which is run at boot and starts all necessary parts of the Operating System (OS) and then spawns the root shell. The service manager ensures that all desired daemons and processes are started at the correct times and remain running, restarting them if they die.
A shell	A (POSIX-compliant) shell **or** a port of a simple shell like Dash[0]. If I decide that implementing my own POSIX-compliant shell is too large of a task, then I will port the source code for Dash to run on my OS, implementing the required syscalls to get it working. This shell will be the main way users interact with the OS.
Documentation	This package contains several parts. I will create documentation for how each part of the OS works, and a guide for how new programmers can get started developing programs for the system. The documentation will be completed in stages as I develop the OS, and I will go back to keep it up-to-date as I make changes to components and add new components.
Multithreading support and multi-core scheduling	The RPi 3 has 4 Central Processing Unit (CPU) cores. This work package will enable users to take advantage of the additional processing power of the other 3 cores for their programs. This includes a threading library similar to `pthread` on Linux systems, and improvements to the scheduler so that it can give different CPU cores to multiple threads owned by the same process. **This work package is optional**.

Table 4: The work packages of the project and their deliverables.
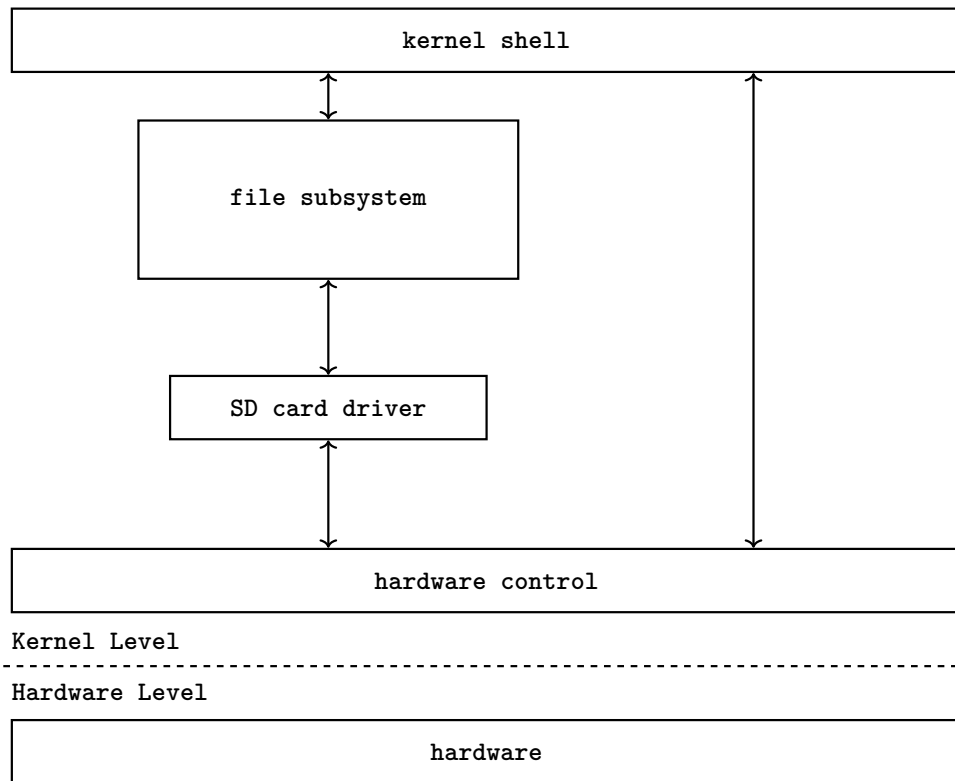
B   Extra diagrams

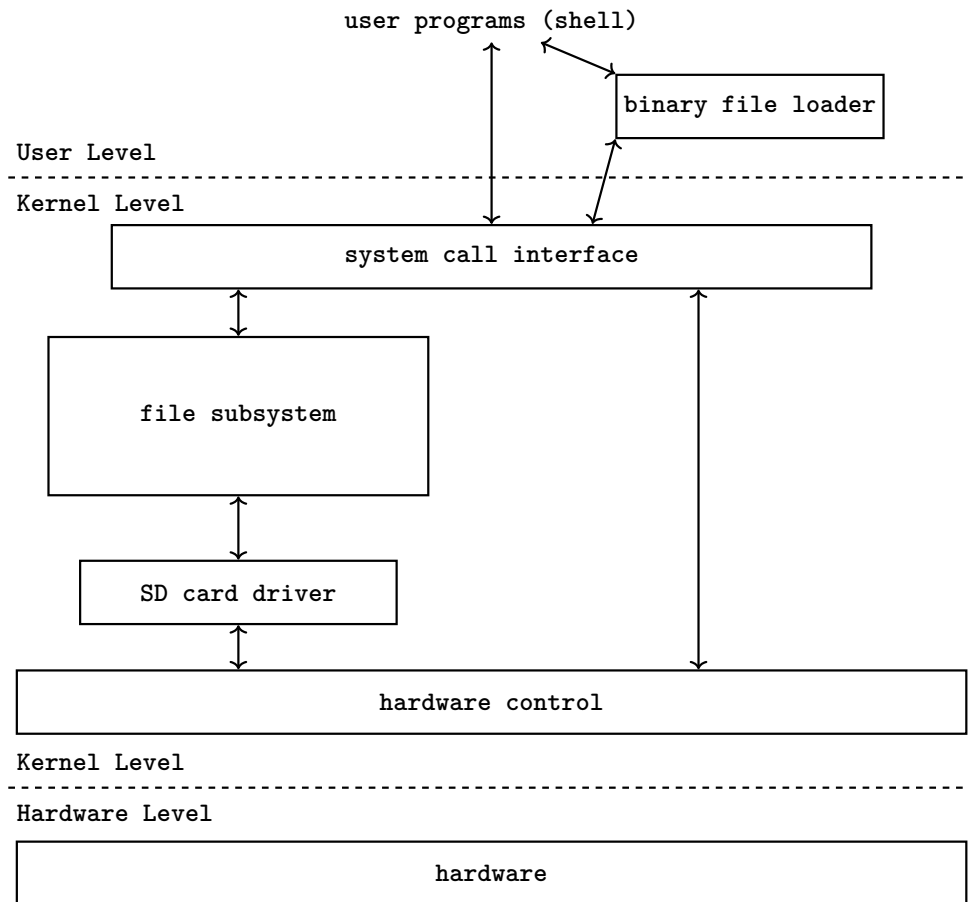Figure 7: The current structure of the Operating System (OS).

user programs (shell)

binary file loader

User Level

Kernel Level

system call interface

file subsystem

SD card driver

hardware control

Kernel Level

Hardware Level

hardware

Figure 8: The 'next-step' for the OS if it was aiming to become Unix after the end of this project.

Figure 9: The structure of the Unix OS.

C Shell future ideas

Here are some 'honorable mentions' which didn't make it into the main report.

- Support for arrow-key navigation (and other special non-character keys) in the console input line
  - History of previous lines entered into the console with up-arrow to get the previous line
  - Cursor position on the line (cursor is between two characters)
  - Backspace to delete one character back, delete to delete one character forwards
- Text pager, like `less` or `more`

Acronyms

arm Arm Microprocessors Ltd.

CPU Central Processing Unit

EL Exception Level

ELF Executable and Linkable Format

FAT File Allocation Table

FS Filesystem

GDB The GNU Project Debugger

GPU Graphics Processing Unit

IO Input / Output

OS Operating System

RPi Raspberry Pi

SBC Single Board Computer

SOC System-On-a-Chip

UART Universal Asynchronous Receiver-Transmitter